


M00: The Agent Lifecycle — See the Whole Picture First

Before you write a single line of code, understand what you're building, why it matters, and where every module fits. This is your map for the entire course.

The Evolution: From Rule-Based AI to Agentic AI

Before you build an agent, you need 75 years of history — in 90 seconds. Agents didn't materialize from thin air in 2024. They're the sixth wave in a long arc that runs from Claude Shannon's information theory in 1948 to today's tool-using, loop-running, memory-keeping systems. **Each era adds a capability without erasing the previous one.** Knowing the whole arc is what keeps you from cargo-culting the latest pattern and lets you reason about *what to use, where, and why* — the exact framing that earns credibility with senior engineers.

SEVEN ERAS OF AI — EACH ERA ADDS, NOTHING IS REMOVED



The Seven Eras — Each With a UCC Example and a Talking Point

Same business problem — *predicting which UCC filings will go delinquent* — solved differently in every era. Each era ends with a one-paragraph **Talking Point** you can memorize verbatim and deliver out loud in twenty seconds.

Era 1 — Foundations & Rule-Based AI (1948 – 2000s)

Key milestones: 1948 — Claude Shannon publishes *A Mathematical Theory of Communication*, the foundation of information theory. 1950 — Alan Turing publishes *Computing Machinery and Intelligence* and proposes the Turing Test. 1955 — John McCarthy coins the term “Artificial Intelligence.” 1997 — IBM Deep Blue beats Kasparov using brute-force search, not learning. 1990s – 2000s — expert systems with hardcoded IF/THEN rules dominate enterprise AI.

What it could do: Follow rules humans wrote. Period.

UCC example: 500+ handwritten rules in a giant decision tree: `IF filing_type=UCC1 AND status=ACTIVE AND days_to_lapse<90 THEN risk=HIGH`. Every state that changed its filing format broke the rulebook.

TALKING POINT — ERA 1

“Rule-based systems were brittle — every edge case required a new rule. They couldn't handle ambiguity, natural language, or anything the programmer didn't explicitly code. They worked for narrow, bounded problems like chess search and expert diagnosis — and broke everywhere else.”

Era 2 — Machine Learning (2000s – 2015)

Key milestones: 2001 — Leo Breiman publishes Random Forests, popularizing ensemble learning. 2006 — Geoffrey Hinton's deep belief networks revive neural networks. 2012 — AlexNet wins ImageNet, kicking off the deep-learning vision era. 2014 — Ian Goodfellow publishes GANs (generative adversarial networks).

What it could do: Learn patterns from data. Classification, prediction, clustering.

UCC example: Train a Random Forest on 10K historical filings to predict delinquency — the *exact* pickle model you'll use in the Prelude below. Input: 6 numbers. Output: a probability.

TALKING POINT — ERA 2

“ML models learn patterns from data, but each model does ONE task. You need separate models for classification, prediction, NER, etc. They can't reason, can't explain decisions, and need structured features — so they can't read raw collateral descriptions or hold a conversation about the result.”

Era 3 — Transformers & NLP Revolution (2017 – 2020)

Key milestones: 2017 — Vaswani et al. at Google publish *Attention Is All You Need*, introducing the Transformer architecture — the single paper that changed everything. 2018 — GPT-1 (OpenAI, Alec Radford) and BERT (Google, Jacob Devlin) ship. 2019 — GPT-2 (1.5B parameters) is initially withheld over misuse concerns. 2020 — GPT-3 (175B parameters) demonstrates few-shot learning across tasks.

What it could do: Process unstructured text. Understand context. Generate coherent language.

UCC example: NLP model reads collateral text (“*All inventory, equipment, and accounts receivable now owned or hereafter acquired*”) and classifies it into categories — inventory, equipment, receivables — with no hand-built dictionary.

TALKING POINT — ERA 3

“Transformers solved the context problem — BERT understood words in context, GPT generated coherent text. But they were still one-task models. GPT-3 changed that by showing one model could do many tasks via prompting alone — that's the discovery the next era was built on.”

Era 4 — Generative AI Explosion (2020 – 2023)

Key milestones: 2020 — GPT-3. 2021 — DALL-E creates images from text; GitHub Copilot writes code alongside developers. 2022 — Stable Diffusion (open-source text-to-image) and Midjourney bring AI art mainstream. Nov 2022 — ChatGPT reaches **100M users in 2 months**, the fastest-growing consumer app ever. Feb 2023 — Claude 1.0 (Anthropic, Constitutional AI). Mar 2023 — GPT-4 (multimodal). Jul 2023 — Claude 2 (100K context); Llama 2 (Meta, open-source).

What changed: AI shifted from UNDERSTANDING content to GENERATING it. Not just classifying a filing — writing a new risk memo. Not just detecting faces — creating photorealistic images from text.

THE FOUR TYPES OF GENERATIVE AI

- **Text generation** — GPT-3/4, Claude, Llama. Articles, emails, analysis.
- **Image generation** — DALL-E, Stable Diffusion, Midjourney. Pictures from descriptions.
- **Code generation** — Copilot, Claude Code, Cursor. Writes and edits software.
- **Audio & video generation** — Sora (OpenAI), ElevenLabs. Speech, music, video from text.

UCC example: Claude reads a 12-page filing and GENERATES a summary, answers questions about it, drafts a risk memo, translates legal language to plain English — all from one model, not retrieved from templates.

THE KEY LIMITATION — the gap agents fill

Generative AI produces content but **cannot TAKE ACTIONS**. Claude can write a beautiful risk report but cannot search the database, check filing status, or call the ML model. *It generates — it doesn't act*. That gap is exactly what agentic AI fills.

TALKING POINT — ERA 4

“Generative AI was the breakthrough that made AI useful to everyone — ChatGPT hitting 100M users in two months proves that. But generative models are REACTIVE: they respond to a prompt. They can't search databases, call APIs, or make decisions in a loop. That limitation is exactly what agentic AI solves.”

Era 5 — Large Language Models Mature (2023 – 2024)

Key milestones: Mid-2023 — enterprise adoption accelerates. Late 2023 — fine-tuning and RAG become standard patterns for domain-specific AI. Early 2024 — Claude 3 family (Haiku / Sonnet / Opus) ships with tiered models for different cost/complexity. Mid-2024 — Claude 3.5 Sonnet becomes the workhorse instruction-follower. Late 2024 — OpenAI o1 introduces explicit chain-of-thought reasoning. 2024 — focus shifts from parameter counts to *capabilities*: multimodality, RAG, tool use.

What changed: LLMs went from demos to production. Companies moved from “let's try GPT” to “let's build products on Claude.” RAG solved the hallucination problem on private data; tiered models solved the cost problem.

UCC example: A bank deploys Claude with RAG over their UCC filing documentation. Analysts ask questions and get answers grounded in *their* data — not hallucinated. But each question is still one prompt → one response. No tool use, no loops, no agent.

TALKING POINT — ERA 5

“2024 was when LLMs went from experiments to production. RAG solved the hallucination problem for domain data. Tiered models — Haiku, Sonnet, Opus — solved the cost problem. But the limitation stayed the same: LLMs respond, they don't ACT. That's what changed with agentic AI.”

Era 6 — Agentic AI (2024 – present)

Key milestones: Early 2024 — Claude tool use API: Claude can call developer-defined functions. Oct 2024 — Claude Computer Use: Claude controls desktop GUIs. Nov 2024 — **Model Context Protocol (MCP)** from Anthropic, the open “USB-C for AI” standard for LLM-tool integration. Late 2024 — OpenAI o1 reasoning models. Early 2025 — **Claude Agent SDK** brings declarative agent building with hooks and sessions. Apr 2025 — **Google Agent2Agent (A2A)** protocol for agent-to-agent communication. May 2025 — **AWS Strands Agents**, an open-source, model-agnostic agent framework. 2025 — agentic-AI startups raise **\$500M+**.

What changed: LLM + tools + loop + memory = an agent that REASONS and ACTS. The critical shift: the LLM goes from *responding* to *deciding* — which tool to call, what to search for, when it has enough information, and when to stop.

UCC example: Agent searches filings across 50 states, discovers name variations by reasoning, calls the ML model, drills into the riskiest lien, writes a narrative report — **the agent you'll build in this course**.

TALKING POINT — ERA 6

“Agentic AI is the convergence of five capabilities that matured at the same time: reliable tool-use APIs, structured output guarantees, large context windows (200K tokens), fast inference (2–5 seconds), and affordable cost (Haiku at \$0.00025/1K tokens). None of these existed in 2022. All of them exist now. That's why agents are possible today.”

Era 7 — The Frontier (2025 – 2026+)

Key milestones: Claude Opus 4.6, GPT-5, Gemini 3 ship with native reasoning baked into the model. Multi-modal agents combine text + vision + audio in a single loop. Agent protocols mature: MCP for tool access, A2A and ACP for agent-to-agent communication (merged under the Linux Foundation). **Distributed “Agentic Intranets”** emerge — agents collaborating across enterprise APIs using natural language. By 2026, organizations begin treating agents as part of the workforce structure: assigning responsibilities, defining ownership.

Where it's going:

- **Phase 1 (2024 – 2025):** Agentic assistants — structured reasoning, planning, tool use inside defined workflows.
- **Phase 2 (2025 – 2026):** Agentic Intranets — agents collaborate ACROSS APIs and enterprise systems.
- **Phase 3 (2026+):** Autonomous orchestration — agents that design, build, and manage other agents.

UCC example: Agent drives the secretary-of-state portal directly, OCRs scanned filings, and hands sub-tasks to a sibling agent that owns “collateral classification” in another business unit.

TALKING POINT — ERA 7

“We're in Phase 1 transitioning to Phase 2. Agents work well within defined workflows; the next step is agents collaborating ACROSS systems. **Two protocols matter:** MCP standardizes agent-to-tool (Anthropic, 2024); A2A standardizes agent-to-agent across orgs (Google → Linux Foundation, 2025). The real shift is that enterprises are starting to treat agents as part of the org chart, not just as tools.”

Deeper coverage in [M24 — The Agent Frontier](#): A2A Agent Cards, Claude Skills, agentic memory frameworks, long-horizon coding agents.

Why Agents Are Possible NOW — the Five Convergences (2022 vs 2026)

The headline is not “the model got smarter.” The headline is the **scaffolding around the model**: function calling, structured output, large context, fast inference, and a **240× cost reduction** since 2022. Without all five shifts below, an agent loop simply doesn't close in production.

Capability	2022	2026	Why It Matters
Tool use API	Did not exist	Native in Claude, GPT, Gemini	Agent can call functions reliably with typed inputs
Structured output	Unreliable prompt-based JSON	Guaranteed via <code>tool_use</code> schema	Agent returns parseable, validated data
Context window	4K – 8K tokens	200K (Claude), 1M+ (Gemini)	Agent holds long conversations + many tool results
Inference speed	10 – 30 seconds	1 – 3 seconds per turn	Multi-turn agent loop completes in reasonable time
Cost per token	\$0.06 / 1K (GPT-3)	\$0.00025 / 1K (Haiku)	240× cheaper — agent loops are finally affordable

“Five years ago you could build a chatbot. Today you can build an agent. The difference isn't a smarter model — it's the tools, context, speed, and cost around the model.”

Market Reality — Data Points Worth Memorizing

If a peer asks “what evidence do you have that agentic AI is real, not hype?” — here are the numbers to keep in your back pocket.

Stat	Where it lands in a conversation
\$500M+ raised by agentic-AI startups in early 2024	Capital is flowing into workflow automation, agent safety, and enterprise integration — not just demos.
20 – 30% operational cost reduction reported by enterprises deploying agentic AI	This is the budget line that gets agent projects funded internally.
35% faster decision automation	Latency to a decision is the metric ops teams actually optimize for.
30 – 50% process throughput improvement with agentic workflows	End-to-end — not just “the LLM step is faster.”
24 of 30 major AI agents launched or majorly updated in 2024 – 2025 (MIT AI Agent Index)	The space is consolidating around a known list of production agents; you can name a few from memory.
Papers mentioning “AI agent” in 2024 exceed all prior years combined	Research output is the leading indicator that industry adoption is about to follow.

THE KEY INSIGHT — capabilities stack, they don't replace

Every era ADDED a capability without erasing the previous ones. When you build a production agent, **all five layers are still in the room**:

- **Rules** live on in guardrails, validation, and policy checks (M16, M17).
- **ML models** live on as tools the agent calls — the Prelude's pickle file is literally one such tool.
- **NLP / embeddings** live on as the engine inside RAG retrieval (M09).
- **LLMs** live on as the agent's brain — Claude IS the reasoner inside the loop.
- **Agents** are the orchestration layer that ties all of the above into a system that *reasons and acts*.

You are not replacing your ML pipeline. You are adding Layer 3 (intelligence + orchestration) on top of everything that already works.

THE 60-SECOND TALKING POINT — the whole arc in one breath

“Generative AI evolved in waves. **Transformers in 2017** enabled understanding language. **GPT-3 in 2020** enabled generating language. **ChatGPT in 2022** brought it mainstream. But LLMs alone are chatbots — they respond, they don't ACT.”

“**Agentic AI is the 2024 wave** — LLMs that use tools, make decisions in a loop, maintain memory, and take actions. The enablers are function-calling APIs, structured output, 200K context windows, fast inference, and a 240× cost reduction since 2022.”

“I'm building agents that combine **ML models** (for prediction), **RAG** (for knowledge), **tools** (for action), **guardrails** (for safety), and **observability** (for production monitoring). That's the full stack, and that's what this course teaches.”

Learning Objectives

By the end of this module, you will be able to:

- ☐ Explain the difference between a chatbot and an AI agent in one sentence
- ☐ Trace the 9-step flow of a real agent interaction from user question to final answer
- ☐ Name the 7 building blocks of a production agent and map each to a course track
- ☐ Describe the 5 stages of the agent lifecycle (Design, Build, Protect, Observe, Deploy)
- ☐ Explain how the same agent patterns built this course — and how you'll learn each one

Prelude: From ML Model to AI Agent

If you've trained ML models or shipped prediction APIs, this section shows what an agent actually *adds* on top. We take one real business problem — UCC delinquency prediction for **Acme Corporation** — and solve it three ways. The data is identical. The pickle file is identical. Only the wrapper changes — and the wrapper changes everything.

Part 1 — The Same Problem, Three Ways

The business question: "Is Acme Corporation likely to become delinquent on secured loans in the next 12 months?" Six features describe the risk: active filing count, state count, collateral types, filing age, amendment frequency, and months to earliest lapse.

Approach 1 — Traditional ML script. A data scientist pickles a RandomForest classifier and exposes `predict_delinquency(features)`. *You* compute the six numbers manually, hand them to the function, get back a probability and label. Fast (milliseconds), reproducible — and totally inert. No data fetch, no explanation, no follow-up.

Approach 2 — FastAPI wrapper. An ML engineer wraps the same pickle in a REST endpoint. POST `{"company_name": "Acme Corporation"}`, the server runs a hardcoded query, computes features, returns rigid JSON. Better — auto-fetches data, validates input. Still inflexible: the query searches `ILIKE 'Acme Corporation'` and misses the filings under `ACME CORP` and `ACME CORP DBA ROADRUNNER SUPPLIES`. Output is still a number.

Approach 3 — Claude agent. The same pickle is now ONE TOOL among three: `search_filings`, `predict_delinquency`, `get_filing_details`. The agent *reasons*: "Try the exact name. Try abbreviations. Try DBAs." It compiles statistics, calls the ML model, drills into the riskiest specific filing, and writes a narrative report citing actual filing numbers. The ML model didn't go away — the agent uses it.

Aspect	Approach 1: Script	Approach 2: FastAPI	Approach 3: Agent
Input	You prepare 6 numbers	Company name (auto-fetch)	Natural-language question
Output	Probability + label	JSON with score + count	Narrative risk report
Name variations	Not handled	Hardcoded ILIKE	Discovered by reasoning
Explanation	None	None	Cites specific filings
Follow-up question	Write new code	Build new endpoint	Just ask
"What if?" scenarios	Retrain model	Not supported	Agent reasons about it
Role of ML model	IS the solution	IS the solution + API	One tool the agent uses
Cost per query	~\$0 (local CPU)	~\$0 (local CPU)	~\$0.01–0.05 (LLM)


The Key Insight

The ML model doesn't disappear when you build an agent — the agent *uses* it. In Approach 1 the model *IS* the product. In Approach 3 the model is one tool the agent calls when it needs a probability. The agent surrounds the model with reasoning, search, and narrative — making it **MORE** useful, not less. Every ML model your team has shipped can become an agent tool tomorrow.

Part 3 — The Three Lanes Visualized

The same input enters all three pipelines. The agent's lane is longer, branchier, and ends with something a credit committee can actually read. Press play to watch each lane fill in.

ANIMATION: SAME DATA, THREE PIPELINES



Same problem · same pickle · three wrappers — only the rightmost lane reads like a memo.

What This Means for the Rest of the Course

You've now seen what an agent *adds*. The remaining 28 modules teach you to build the rightmost lane from scratch: the loop (M03–M05), the tools (M05–M07), the search (M09 RAG), the reasoning patterns (M10–M12), the guardrails that prevent runaway loops (M16–M18), the observability that lets you trace every turn (M19–M20), and the deployment paths to local Docker, Cloud Run, and Lambda (M21–M22). Every ML model your team owns can become a tool — and every tool makes the agent more capable.

Why Agents: The Business Case

You just saw three approaches solve the same UCC problem. A sharp reader will notice something: in Approach 2 the ML model is wrapped in FastAPI — and in production the agent (Approach 3) will *also* be wrapped in FastAPI (M22B teaches exactly that). So what is actually different? This section answers that question, then lays out the seven concrete benefits an agent provides, and finally tells you when *not* to use one.

Wait — Both End Up in FastAPI. So What's the Difference?

The infrastructure is identical. What changes is **where the decision logic lives**. Compare the two server bodies side by side:

ARCHITECTURE COMPARISON — SAME INFRASTRUCTURE, DIFFERENT BRAIN

Approach 2 — ML in FastAPI

POST /predict
{ "company_name": "Acme" }

↓

1. Parse input

2. Query DB ← hardcoded SQL

3. Load pickle

4. Predict

5. Return JSON

Logic: YOUR CODE decides everything

→ { "prediction": "HIGH RISK", "probability": 0.823 }

Approach 3 — Agent in FastAPI

POST /query
{ "question": "Assess risk for Acme" }

↓

1. Parse input

2. Call Claude

3. Claude THINKS ← what to search?

4. Claude calls YOUR tools

5. Claude loops & synthesizes

6. Claude writes narrative

Logic: CLAUDE decides what to do

→ "Acme has 8 active filings across 4 states. ML predicts HIGH RISK (82.3%) primarily because filing CA-2024-001 covers all assets and lapses in 8 months..."

What's the same: FastAPI server, an HTTP endpoint, Docker container, auth, rate limits, deployment to Cloud Run or Lambda. **What's different:** who decides what to do once a request arrives.

Aspect	ML Model in FastAPI	Agent in FastAPI
Who decides what to query?	Your hardcoded SQL	Claude reasons about what to search
Who handles name variations?	Your <code>ILIKE</code> pattern	Claude discovers them
Who picks which data to look at?	Your code, fixed order	Claude, based on what it finds
Who formats the response?	Your template	Claude writes natural language
What changes when logic changes?	Your code + redeploy	The system prompt (no redeploy)

The Three-Layer Stack — What Agents Actually Add

The cleanest way to see the difference is as three layers. Two of them are identical in both approaches. The third is what an agent adds.

ANIMATED: THE INTELLIGENCE LAYER APPEARS

↓ what agents ADD on top

LAYER 3 — INTELLIGENCE (Claude)

NEW with agents

reasoning · planning · synthesis · explanation

LAYER 2 — CAPABILITIES

same in both

search_filings() · predict_delinquency() · ML model

LAYER 1 — INFRASTRUCTURE

same in both

FastAPI · Docker · HTTP · auth · rate limits

Layer 1 (Infrastructure) and Layer 2 (Capabilities, including the ML model) are identical in both approaches. The ML model *does not move* — it stays in Layer 2 in both. The difference is what sits *above* it: in Approach 2 your hand-written if/else logic; in Approach 3 Claude's reasoning. **Agents don't replace ML models — they put a reasoning layer on top of them.**

The Career Insight

In five years most APIs will still be FastAPI (or equivalent). Most ML models will still be pickle, ONNX, or TensorFlow files. The change is Layer 3 — the reasoning that decides how to use the tools and models. That is what this course teaches you to build.

The Cost-Benefit Reality

An agent costs more per request and is slower. It also collapses development time and produces output a non-technical user can read.

Metric	ML in FastAPI	Agent in FastAPI
Response time	50–200 ms	3–15 seconds
Cost per request	~\$0	\$0.003–0.075
Development time for v1	2–3 days	4–6 hours
Time to add a new question type	1–2 days (new endpoint)	0 (Claude handles it)
Time to handle a new edge case	Hours (find, code, test, deploy)	0 (Claude reasons about it)
Maintenance burden	High (every change = code)	Low (tools rarely change)
Explainability	Manual feature importance	Built-in narrative
User training	API documentation	None — natural language

The 7 Benefits — UCC Examples

1. Reasoning replaces hardcoded logic. The Approach 2 script searched five hardcoded name variants. The agent discovered **ACME CORP DBA ROADRUNNER SUPPLIES** by reasoning "*let me also check for DBAs*" — a step nobody programmed. A real credit-risk team replaced 200+ name-variation rules with an agent and saw match rate jump from 78% to 94%.

2. **Natural language in, structured action out.** Traditional: `POST /predict {"company_name": "ACME CORPORATION", "state": "NY"}`. Agent: "What's the lien exposure for Acme across the northeast?" Adoption widens from five engineers who know the API to fifty analysts who can just ask.
3. **Explainability is built in.** ML model says `0.823`. Agent says "HIGH (82.3%) primarily because filing CA-2024-001 covers all assets and lapses in 8 months — if it lapses without renewal, \$2.4M of collateral becomes unsecured." OCC and FDIC examiners can read the second; they cannot read the first.
4. **Follow-up questions without new code.** Traditional path requires `/predict`, then `/predict-by-state`, then `/compare-entities`, then `/what-if-continuation` — four endpoints. The agent handles all four naturally with the same three tools. New question types ship at zero engineering cost.
5. **Multi-source synthesis.** The prelude agent searched filings (Tool 1), ran the ML model (Tool 2), drilled into the riskiest filing (Tool 3), and wove the three into one report. Writing that synthesis as a script is 200+ lines and breaks when a fourth data source arrives. Claude handles the joining.
6. **Graceful handling of incomplete data.** Search returns nothing for Texas. A script silently omits TX. The agent says: "Searched TX — no active filings found. Last filing TX-2021-005 was terminated in 2022." Analysts stop asking "did you check Texas?" because the agent volunteers what it checked.
7. **The ML model gets smarter context (not replaced).** The model still sees six numbers. The agent wraps those numbers with which filings produced them, what the collateral descriptions say, and what a continuation or termination would change. The model gives the probability; the agent gives the story. **Agents don't replace ML models — they make them more useful.**

When NOT to Use Agents

Agents are not the right answer for every problem. Reach for a script, an API, or a rule engine when:

Situation	Why Not an Agent
Batch processing 1M records	Agent cost: \$10K+. Script cost: ~\$0.
Sub-100ms response required	Agent latency: 3–15 s. Script: milliseconds.
Deterministic compliance check	Must be reproducible bit-for-bit. Agents are non-deterministic.
Simple CRUD operations	No reasoning needed. Over-engineering.
No human will read the output	If the next consumer is another system, a structured API is better.

The Decision Rule

If a **human** needs to UNDERSTAND the output, consider an agent. If a **machine** needs to CONSUME the output, use an API. The agent costs \$0.015 and 10 extra seconds — cheap if the alternative is a human analyst spending 30 minutes at \$50–100/hour, expensive if the next caller is a downstream batch job that just needs a number.

What Is and Is Not an Agent — The Clear Boundary

Students hear "agent" and assume any code that calls Claude is one. It is not. There are **three distinct levels** of LLM-powered programs, and only the third is actually an agent. The boundary between them comes down to a single question: *who decides what happens next — your code, or Claude?*

The Three Levels at a Glance

- Level 1: LLM Call** — your code asks, Claude answers, done.
 - Level 2: LLM Workflow** — your code orchestrates a fixed sequence of LLM calls.
 - Level 3: Agent** — Claude decides which tool to call and when to stop.
- In Levels 1 and 2 **your code makes every decision** (what to ask, in what order, when to stop). In Level 3 **Claude makes the decisions** based on what it discovers at runtime. That single shift — from your `if/else` to Claude's reasoning — is the entire boundary.

Level 1: LLM Call (NOT an agent)

One request, one response. No tools. No loop. Your code decided to call the LLM exactly once and end.

```
# level1_call.py — a single Claude call
import anthropic
client = anthropic.Anthropic()
```

```

response = client.messages.create(
    model="claude-sonnet-4-20250514",
    max_tokens=1024,
    messages=[{"role": "user",
                "content": f"Summarize this UCC filing: {filing_text}"}]
)
print(response.content[0].text)
# Done. One call. One response. YOUR code decided to summarize.

```

Real Level 1 examples: chatbot answering a single question, text summarizer, translation service, code reviewer, email draft generator.

Level 2: LLM Workflow (NOT an agent)

Multiple LLM calls in a *fixed* sequence. Each step's output feeds the next. Your code decided every step before runtime.

```

# level2_workflow.py – three LLM calls in a fixed order
# Step 1: extract (YOUR code chose this as step 1)
extract = client.messages.create(model="claude-sonnet-4-20250514", max_tokens=1024,
    messages=[{"role": "user",
                "content": f"Extract debtor + secured party from: {filing_text}"}])
entities = extract.content[0].text

# Step 2: classify (YOUR code chose this as step 2)
classify = client.messages.create(model="claude-sonnet-4-20250514", max_tokens=1024,
    messages=[{"role": "user",
                "content": f"Classify risk for these entities: {entities}"}])
risk = classify.content[0].text

# Step 3: report (YOUR code chose this as the final step)
report = client.messages.create(model="claude-sonnet-4-20250514", max_tokens=2048,
    messages=[{"role": "user",
                "content": f"Write a risk report. Entities: {entities}. Risk: {risk}"}])
print(report.content[0].text)
# Three calls. Always extract → classify → report. YOUR code chose the order.

```

Real Level 2 examples: ETL pipelines (extract → transform → load), content pipelines (research → draft → edit → publish), CI/CD review flows. The tell: if you draw the program as a flowchart, every path is known *before* runtime — there are no decision diamonds where Claude picks the next step.

Level 3: Agent (THIS is an agent)

A loop where **Claude** picks the tool, the arguments, the order, and the stopping point. The execution path is unknown until runtime.

```

# level3_agent.py – Claude decides what to do next
tools = [
    {"name": "search_filings", "description": "Search UCC filings by name", "input_schema": {...}},
    {"name": "get_details", "description": "Get filing details", "input_schema": {...}},
    {"name": "calculate_risk", "description": "Calculate risk score", "input_schema": {...}},
]
messages = [{"role": "user",
              "content": "What is the lien exposure for Acme Corporation?"}]

while True:
    response = client.messages.create(
        model="claude-sonnet-4-20250514",
        max_tokens=4096,
        tools=tools,
        messages=messages,
    )

    if response.stop_reason == "end_turn":
        print(response.content[0].text)
        break # Claude decided it has enough info

# Claude chose the tool – not your code
for block in response.content:
    if block.type == "tool_use":
        result = execute_tool(block.name, block.input)
        messages.append({"role": "assistant", "content": response.content})
        messages.append({"role": "user", "content": [

```

```
        {"type": "tool_result", "tool_use_id": block.id, "content": str(result)}
    })
    # Loop – Claude decides what's next based on the result
```

Three things must all be true for Level 3: *tools* (Claude can act on the world), *a loop* (Claude can act again), *decisions* (Claude chooses based on what it found). Remove any one of the three and you drop back to Level 1 or Level 2.

The Decision Matrix


One row per dimension that actually distinguishes the three levels. If you can answer these seven questions about your program, you can label it instantly.

Question	Level 1: Call	Level 2: Workflow	Level 3: Agent
How many LLM calls?	1	2+ (fixed count)	Unknown (dynamic)
Who decides the sequence?	YOUR code	YOUR code	Claude
Uses tools?	No	No (or fixed tools, fixed order)	Yes (Claude picks)
Has a loop?	No	No (linear pipeline)	Yes (until done)
Adapts to results at runtime?	No	No	Yes
Path known upfront?	Yes	Yes	No
Checks <code>stop_reason</code> ?	No	No	Yes (<code>tool_use</code> vs <code>end_turn</code>)

The Litmus Test

If you replace the LLM with hardcoded responses and the program still works the same way — it is NOT an agent.

For Level 1 and Level 2, swap Claude with a canned string and the program runs identically. The LLM is just a function that returns text. For Level 3, the agent's behavior depends on Claude's runtime *decisions* — which tool, with which arguments, when to stop. A canned response cannot make those decisions, so the agent breaks. That fragility under hardcoding is the proof that Claude is doing the deciding — and that proof is the agent.



Common Gray Areas

"My code calls Claude, then based on the answer calls Claude again. Agent?" — No. Your `if/else` branched, not Claude. That is Level 2 — a workflow with a conditional edge.

"My code calls Claude with tools, Claude uses one tool, I return the result. Agent?" — Borderline. If it always fires exactly one tool and stops, it is Level 1 with a tool attached — closer to a router than an agent. If you loop until `stop_reason == "end_turn"` and Claude is allowed to call multiple tools across multiple turns, it is Level 3.

"I use LangChain's AgentExecutor (or any framework). Is it an agent?" — Yes, *if* Claude has tools, a loop, and is making the decisions about what to call next. The framework does not change the definition; the pattern does. A "framework agent" wired as a fixed pipeline is still Level 2.

"My code asks Claude to choose a function, then my code calls it directly. Agent?" — Closer to Level 2.5 — the "router pattern". Claude makes one routing decision per request but does not loop or adapt to results. Useful, but not yet an agent. We cover routers explicitly in M13.

Now you have the boundary, sharp and operational. Everything before **M05** in this course is Level 1. **M05 adds tools. M12 adds the loop and the `stop_reason` check.** That is the precise moment you cross from program to agent — and the rest of the course is about doing it well.

What Is an AI Agent? (And What Isn't One?)

You've probably used ChatGPT, Claude, or another AI chatbot before. You type a question, it gives you an answer, and the conversation is over. That's a chatbot — and it's useful, but it has a fundamental limitation: it can only *talk*.

An AI agent is different. An agent doesn't just talk — it **acts**. It can search databases, call APIs, and read files. It makes decisions based on what it finds. And then it loops back to do more work until the job is actually done. The key word is *loop*: an agent keeps going until it has a complete answer, not just a single response.

Everyday Analogy

Before: Imagine you walk up to an information desk at an airport. You ask "When does the next flight to Chicago leave?" The person behind the desk answers from memory: "I think it's around 3pm." That's a chatbot — one question, one answer from what it already knows.

The pain: But what if the answer is wrong? What if there are multiple flights, and you need the cheapest one? The information desk can only tell you what they remember — they can't look anything up, check real-time pricing, or book a ticket for you.

The agent version: Now imagine a travel assistant who hears your question, pulls up the flight database, checks 4 different options, compares prices, notices you have frequent flyer miles on United, applies them, and books the best option — all from a single request. That's an agent. Same question, but the assistant *used tools, made decisions, and kept working* until the task was complete.

What this looks like in real code: When Claude decides to use a tool, it doesn't magically connect to a database. It sends back a structured JSON message like this: `{"type": "tool_use", "name": "search_flights", "input": {"destination": "Chicago", "date": "2025-03-15"}}`. Your code receives that message, runs the actual search, and sends the results back. That JSON handshake is the heart of every agent interaction.

Technical Definition

An AI agent is a program that combines three capabilities: (1) an LLM as its reasoning engine (the "brain" that reads situations and decides what to do), (2) **tools** it can call to interact with the outside world (databases, APIs, files — the "hands"), and (3) a **loop** that keeps the agent thinking and acting until the task is done (not just one response, but as many rounds as it takes).

Here's the important nuance: an agent is *not* autonomous AI running on its own. It's not sentient, and it's not making decisions 24/7 without supervision. It's **your code** that uses an LLM to make decisions inside a controlled loop. You define what tools it can use, what it's allowed to do, and when it should stop. The LLM provides the reasoning; your code provides the structure and guardrails.

Chatbot vs Agent — Side by Side

The animation below shows the same question handled two ways. On the left, a chatbot gives a single response from memory. On the right, an agent thinks, uses tools, and loops until it has a complete answer. Watch how many steps the agent takes compared to the chatbot.

ANIMATION: CHATBOT VS AGENT

CHATBOT (ONE TURN)

What's Acme Corp's total lien exposure?

I don't have access to filing databases. I can explain what a UCC lien is in general terms...

DONE — SINGLE TURN

AGENT (MULTI-STEP LOOP)

What's Acme Corp's total lien exposure?

I need to search the UCC filing database...

`search_filings("Acme Corporation")` → 7 results

Check for name variations too...

`search_filings("ACME CORP")` → 3 more results

`get_risk_profile("acme-id")` → High Risk

Acme Corp has 10 UCC filings across NY, CA, TX totaling \$4.2M in lien exposure. Risk: High.

DONE — 6 STEPS, 3 TOOL CALLS

Why It Matters

This isn't just a theoretical difference. In a real bank's compliance department, an analyst checking lien exposure might spend **45 minutes** manually searching 3 different state databases, cross-referencing name variations, and compiling results into a report. An agent does the same work in **15 seconds** — and catches the name variation the human might miss. That's the promise: agents don't just answer questions, they *do work*.

Common Misconceptions About AI Agents

"Agents are autonomous AI that run on their own, right?" — No. An agent runs when your code calls it and stops when the task is done. It doesn't think, plan, or act between calls. It's a program with an LLM inside a loop — not a self-directed entity.

"An agent is basically a smarter chatbot?" — Not exactly. A chatbot responds from memory in one turn. An agent can use tools, make decisions, and loop multiple times. The difference isn't intelligence — it's the ability to *act* on the world and *iterate* until done.

"If I build an agent, it might go rogue or do things I didn't intend?" — Only if you don't add guardrails. The agent can only use tools YOU define, and you control the loop. A well-built agent has iteration limits, input validation, output checking, and human approval gates. That's why this course dedicates an entire track (M16-M18) to safety.

"Agents must be complex — hundreds of lines of code?" — The core agent pattern is about 10 lines: a while loop that calls an LLM, checks if it wants to use a tool, runs the tool, and repeats. Everything else (memory, planning, guardrails) is added incrementally. You'll have a working agent by Module 5.


Script vs Agent — Why This Course Exists

Before we show you a full agent in action, let's answer the obvious question: *why not just write a regular Python script?* Consider the task from the previous section — finding total lien exposure for Acme Corporation. Here's how a script and an agent approach the same problem:

The Core Difference

Aspect	🔧 Script	🤖 Agent
Name variations	Hardcoded list YOU maintain	Discovered by reasoning at runtime
States to search	Fixed list	Dynamic — based on findings
Decision logic	if/else chains YOU write	Claude reasons about what to do next
New edge cases	Code change + redeploy	Handled by reasoning — no code change
Follow-up questions	Build a new function	Natural conversation continuation
Development time	Days (handle every case)	Hours (define tools + loop)

ANIMATION: SCRIPT APPROACH VS AGENT APPROACH



The Key Insight

Why It Matters

An agent is a script that replaced the hardcoded decision logic with an LLM. Instead of YOU writing every if/else and every loop condition, Claude *reasons* about what to do next. Your code provides the **tools** — what the agent *can* do. Claude provides the **logic** — what the agent *should* do. The tool-use loop you'll write (about 15 lines) replaces hundreds of lines of decision logic. But you still write the tools — the agent doesn't magically connect to databases or APIs. You provide the hands. Claude provides the brain.

When Scripts Are Better

Agents are **not** always the answer. A traditional script wins when:

- **The task is 100% deterministic** — no reasoning needed (file format conversion, CSV → JSON)
- **Speed matters more than flexibility** — scripts are 100× faster than an LLM round-trip
- **Cost matters** — every agent turn is an API call, and API calls cost money
- **The logic never changes** — agents add unnecessary complexity to fixed workflows
- **You need guaranteed reproducibility** — same input must produce the exact same output every time

Batch record insertion and cron jobs with fixed logic are script territory. So is data validation against a known schema and simple CRUD operations. This course teaches you to reach for agents when the problem demands reasoning, not as a default for everything.

See an Agent in Action — Live Demo Walkthrough

Let's make this concrete. Imagine you work at a bank, and your job is to assess lien risk for corporate borrowers. Today, a loan officer asks: *"What's the total lien exposure for Acme Corporation across all states?"*

Below is exactly what happens behind the scenes when an agent handles this question. This is the same [UCC Filing Research Agent](#) that we'll use as a running example throughout the course. Every step maps to concepts you'll learn in specific modules.

The 9-Step Agent Interaction



What Just Happened?

The agent made **three key decisions** that a chatbot never could: (1) it decided it needed to search a database (tool selection), (2) it decided to search again with a name variation (autonomous reasoning), and (3) it decided to pull a risk profile for a complete answer (multi-step planning). Each decision came from Claude's LLM reasoning — the same reasoning you'll learn to harness starting in Module 1.

This walkthrough isn't hypothetical — it's the exact architecture you'll build. By Capstone 1 (after Module 7), you'll have a working version of this agent. By Capstone 5 (after Module 22), you'll have a production-grade version with guardrails, monitoring, and deployment.

The Agent Architecture — Building Blocks Map

Every production agent, no matter how simple or complex, is built from the same 7 components. Think of them like the parts of a human body — each has a specific job, and they all work together. The animation below shows each component lighting up and connecting to the course modules where you'll learn to build it.

Everyday Analogy


Before: Think about a doctor seeing a patient. A medical student who only learned anatomy but never practiced diagnosis would be dangerous. A doctor who can diagnose but has no access to lab tests can't confirm anything. A doctor who can diagnose and order tests but never checks the results is negligent.

The pain: A real doctor needs ALL the skills working together — knowledge (brain), tools (lab access), memory (patient history), judgment (diagnosis plan), safety checks (drug interaction alerts), monitoring (follow-up), and practice setting (clinic or hospital).

The mapping: An AI agent works the same way. The LLM is the brain (knowledge + reasoning). Tools are the hands (database access, API calls). Memory is the patient chart (conversation history, RAG). Planning is the diagnosis process (breaking complex problems into steps). Guardrails are safety protocols (preventing harmful actions). Observability is the monitoring system (logs, traces). Deployment is the practice setting (where it runs). Remove any one piece, and the agent is incomplete.

What this looks like in practice: Here's a simplified agent config showing all 7 blocks wired together: `{"brain": "claude-sonnet-4-6", "tools": ["search_filings", "get_risk_profile"], "memory": "conversation_history[]", "plan": "ReAct loop", "guardrails": {"max_turns": 10, "banned_tools": ["delete_filing"]}, "observability": "langfuse_trace_id", "deployment": "cloud_run_endpoint"}`. Each field maps to a building block. By Module 22, you'll understand every one of these.

ANIMATION: THE 7 BUILDING BLOCKS OF A PRODUCTION AGENT



Here's what each component does and why it matters:


Component	What It Does	Without It...	Course Track
Brain (LLM)	Reads input, reasons about it, decides what to do next, generates text	No reasoning — just static rules or pattern matching	Track 1: M01-M04
Tools	Calls APIs, queries databases, reads files, sends emails	Can only respond from training data — no real-time info	Track 2: M05-M07
Memory	Remembers conversation history, retrieves relevant documents (RAG)	Forgets everything between messages — asks "who are you?" every turn	Track 3: M08-M11
Plan	Breaks complex tasks into subtasks, decides execution order	Can only handle simple, one-step requests	Track 4: M12-M15
Guardrails	Validates inputs, checks outputs, escalates to humans when unsure	No safety net — generates harmful content, runs expensive loops	Track 5: M16-M18
Eyes	Logs every decision, traces tool calls, monitors performance	A black box — when it breaks, you have no idea why	Track 6: M19-M20
Home	Runs in production: API design, containerization, scaling, cost control	Works on your laptop, but nobody else can use it	Track 7: M21-M22

"Most tutorials stop at the Brain and Tools — that's only 2 of 7 components. This course covers all seven, because an agent that works on your laptop but can't be trusted in production isn't useful."

The Agent Lifecycle — From Idea to Production

Building an agent isn't just "write code and ship it." Production agents go through 5 stages, each adding a critical layer. The animation below shows these stages flowing left-to-right, with the course tracks mapped to each.

ANIMATION: THE 5 STAGES OF AN AGENT LIFECYCLE



Let's walk through each stage with the UCC Filing Research Agent as our example:

1. **Design** — What should the agent do? What tools does it need? What data does it access? For our UCC agent: it needs to search a filing database, handle name variations, and produce risk assessments. You'll learn how LLMs think (M01-M04) and how to define tools (M05-M07) in this stage.
2. **Build** — Write the code. This means defining tools, engineering prompts, and wiring up the agent loop. You might also add a RAG pipeline for policy documents. This is where the UCC agent gets its `search_filings` and `get_risk_profile` tools wired up. Covered in Tracks 2-4 (M05-M15).
3. **Protect** — What if someone asks the agent to delete filings? What if it hallucinates a risk score? What if it runs 500 tool calls and costs \$100? This stage adds guardrails to validate inputs and outputs. It also adds human-in-the-loop approval for high-stakes actions and cost controls to prevent runaway spending. Track 5 (M16-M18).
4. **Observe** — When the agent produces a wrong answer, how do you figure out why? Tracing every LLM decision, logging tool calls, monitoring accuracy over time. Without this, production agents are black boxes. Track 6 (M19-M20).
5. **Deploy** — Ship it: wrap the agent in an API, containerize it, deploy to cloud, optimize costs. The UCC agent becomes a service that any analyst in the bank can call. Track 7 (M21-M22).

Most Tutorials Stop at Stage 2

Search "build an AI agent" on YouTube and you'll find hundreds of tutorials that cover Design and Build. Almost none cover Protect, Observe, or Deploy. This course covers **all five stages** — because an agent that works in a demo but breaks in production, costs too much to run, or can't be debugged when it fails isn't a real product. It's a prototype.

Common Misconceptions About the Agent Lifecycle

"I can add guardrails later, after the agent works." — Dangerous approach. By the time you discover your agent can hallucinate risk scores or run 200 tool calls per query, you've already deployed it. Build guardrails into your agent loop from the start, even if they're simple (like a max iteration count).

"Observability means just adding print statements." — Print statements disappear when the server restarts. Production observability means structured logs, distributed traces, and dashboards that show you *why* the agent chose a particular tool on a specific request last Tuesday. You'll learn this in Modules 19-20.

"The lifecycle is waterfall — finish one stage before starting the next." — Not at all. You'll iterate between stages constantly. Building a new tool? You also need a guardrail for it and a trace for it. The lifecycle describes what concerns exist, not a strict order of operations.

How Agents Actually Work — The API Reality

Here's a common misconception: many people think an AI agent is a separate application running somewhere, making decisions on its own. The reality is much simpler — and more empowering, because it means *you* are always in control.

An agent is **your code** that calls Claude's **API** in a loop. There's no magic runtime, no background daemon, no persistent process making autonomous choices. Your code sends a request. Claude's servers send a response. Your code decides what to do with that response. Repeat.

This matters because it means the agent only runs when your code tells it to. Between API calls, nothing is happening. The agent doesn't "think" while waiting. It doesn't accumulate knowledge between sessions (unless you explicitly save state). Every call is stateless — you send the full conversation history each time, and Claude responds as if it's seeing the conversation for the first time.

Here's the entire pattern in pseudocode — every agent, from a simple calculator to a multi-million-dollar enterprise system, is a variation of this:

```
// The universal agent pattern
user asks a question

while (claude wants to use a tool):
    send the question (+ history) to claude
    if claude returns a tool request:
        run the tool
        send the result back to claude
    else:
        return claude's answer to the user
```

What Just Happened?

Look at that pseudocode again. The **while** loop is the entire agent. Your code sends a message to Claude, and Claude responds with either a **tool request** ("please call **search_filings** with these inputs") or a **final answer** ("here's your report"). If it's a tool request, you run the tool and send the result back — that's one iteration. If it's a final answer, you return it to the user and the loop ends. That's it. Really.

The rest of this course teaches you every variation, optimization, and production concern around this loop. You'll learn how Claude decides whether to call a tool (Module 5), how to handle multiple tools at once (Module 6), how to give Claude access to your private data (Module 9), how to prevent infinite loops (Module 16), and how to trace every step for debugging (Module 19).

Where Agents Run

During development, your agent runs on your laptop — you type a command, the loop executes, you see the result. This is how you'll work through most of this course.

In production, the same loop runs on a server. Typically it sits behind a REST API or webhook that receives user requests and returns agent responses. A loan officer clicks "Analyze" in a web app, which sends a request to your server, which runs the agent loop, which calls Claude's API, which calls your tools, and returns the final answer. The loan officer never sees the loop — just the result.

Some agents don't have a user at all. They run in CI/CD pipelines (automated code review on every pull request), on cron schedules (daily compliance report generation), or inside chat interfaces like Slack (customer support bots). The loop is the same everywhere — only the trigger changes. You'll learn to deploy agents in all these contexts in Modules 21-22.

How a Claude Agent Built This Course

Before we start teaching you to build agents, let us show you one in action — **the one that built the course you're reading right now.**


This entire course — every module, every animation, every quiz question, every code example — was generated by an AI agent using the exact same patterns you're about to learn. The agent is called Claude Code, and it's Anthropic's own agentic coding tool.

Claude Code is an AI agent that reads files, writes code, executes terminal commands, and iterates until the output meets quality standards. It runs locally on the course author's machine. It uses Claude (the same model you'll use) as its LLM brain. And it follows the *exact same architecture* that you'll build in this course.

The Course-Building Agent Architecture

Here's what the course-building agent looks like. Every component maps to a concept you'll learn:

ANIMATION: THE AGENT THAT BUILT THIS COURSE



Let's map each component to the course modules where you'll learn to build it yourself:

Agent Component	What It Did for This Course	Where You'll Learn It
<code>CLAUDE.md</code>	Project memory — told the agent the design system, quality standards, and depth rules to follow	M25 (Claude Code Mastery)
Slash commands	Predefined workflows: <code>/generate-module</code> , <code>/fix-explanations</code> , <code>/review-module</code>	M25 (Claude Code Mastery)
Prompt files	Loaded design specs, module briefs, and cert tips on-demand — like RAG without a vector database	M09 (RAG), M25 (Claude Code)
Read / Write / Edit tools	Read existing modules for consistency, wrote new HTML files, edited specific sections	M05 (Function Calling)
Bash tool	Ran file checks, counted sections, verified HTML structure	M15 (Code Interpreter)
<code>ReAct</code> loop	Think → Read → Generate → Review → Edit → Repeat until quality passes	M12 (ReAct Pattern)
Quality checklist	16-point validation after every module (guardrails for the agent)	M16-M17 (Guardrails)

But How Does Claude Actually Build the HTML?

This is the question students always ask — and the answer reveals something important about how LLMs work.

Claude doesn't use a website builder, a template engine, or a React framework. There's no WordPress, no static site generator, no drag-and-drop editor. Here's what *actually* happens:

Step 1: Load context (the Read tool). The agent reads specification files from disk — the CSS design system, the 14 depth rules for explanation quality, the module brief describing what concepts to cover. These files go into Claude's `context window` — its working memory for this task. This is functionally identical to the `RAG` pattern you'll learn in Module 9: load relevant knowledge before generating.

Step 2: Generate HTML (the LLM brain). Claude learned HTML, CSS, and JavaScript during training — from millions of web pages, documentation sites, and interactive tutorials. When it generates a module, it combines the design specifications (from Step 1) with its knowledge of web technologies to produce a complete, self-contained HTML file. It writes this character by character, `token` by token — just like it writes any text response. There is no template with blanks to fill.

Step 3: Write the file (the Write tool). The agent saves the generated HTML to disk — about 100-200KB of self-contained code with all CSS and JavaScript inline.

Step 4: Review its own work (Read + Grep tools). The agent reads the file back and checks it against quality rules. Does it have at least 5 quiz questions? Are all sections present? Are there ARIA labels for accessibility? This is the agent's guardrail — reviewing its own output before reporting success.

Step 5: Edit specific sections (the Edit tool). When the human says "make the animation slower," Claude doesn't regenerate the entire 150KB file. It finds the specific CSS value and changes it. Only that line changes. This makes iterations fast (seconds) and precise (no unintended side effects).

The Workflow — Step by Step

ANIMATION: HOW ONE MODULE GETS BUILT

1

Human types `/generate-module M09`

Command

2

Agent reads 8 specification files

Read

3

Agent generates complete HTML with animations, code, quizzes

LLM Brain

4

Agent writes file to disk (100-200KB)

Write

5

Agent runs 16-point quality checklist

Grep + Bash

6

Human previews in browser, requests changes

HITL

7

Agent edits specific sections (not full regeneration)

Edit

8

Repeat steps 5-7 until quality standards met

Loop

The Key Insight

The course-building agent uses **no special capabilities**. It uses the same tools and patterns you'll learn: file reading (M05), file writing (M05), code execution (M15), loading specs before generating (RAG, M09), generate → review → fix loop (ReAct, M12), and quality checks (Guardrails, M16-M17). The "magic" isn't in the tools — it's in the **specifications** (prompt files that define exactly what to generate) and the **iteration loop** (review and fix until quality passes). By Module 22B, you'll be able to build a system like this yourself.

The Numbers

Metric	Value
Time per module (agent work)	~4 minutes
Time per module (human review)	~10-15 minutes
Total modules generated	30 (including capstones)
Total API cost	~\$20-30
Output per module	80-200KB self-contained HTML

Reflection Exercise

This module is intentionally code-free. Before you start building in M01, take a moment to think about *your* agent ideas.

Reflection 1: Your Agent Idea

Think of a repetitive task in your work that involves: **looking up information**, **making decisions based on rules**, and **producing a report or response**. That task is a candidate for an agent. Write it down — by the end of this course, you'll be able to build it.

Examples: reviewing insurance claims against policy criteria, checking inventory levels and reordering, answering customer questions from a knowledge base, analyzing financial filings for risk factors.

Reflection 2: Generate → Review → Fix

The course-building agent follows a simple pattern: generate output, review it against quality rules, fix the issues, repeat. What's a similar **generate** → **review** → **fix** workflow in your domain that an agent could handle?

Examples: drafting legal documents and checking them against templates, generating test cases and verifying they cover edge cases, writing marketing copy and ensuring brand guideline compliance.

Module Summary

You've seen the whole movie. Here's what you now know:

- **An agent = LLM + Tools + Loop.** It doesn't just talk — it thinks, acts, observes, and repeats until the job is done.
- **7 building blocks** make up a production agent: Brain, Tools, Memory, Plan, Guardrails, Eyes (observability), and Home (deployment).
- **5 lifecycle stages:** Design → Build → Protect → Observe → Deploy. Most tutorials stop at Build.
- **The universal agent pattern** is a while loop: call LLM → check for tool use → run tool → repeat.
- **This course was built by an agent** using the same patterns you'll learn — Read → Generate → Review → Edit → Repeat.

What's Next: M01 — The LLM Mental Model

Now you understand the big picture. In Module 1, you'll zoom in on the Brain — the LLM that powers every agent. You'll learn how Claude actually generates text (it's not looking up answers in a database), why "temperature" changes the output, and what tokens are. This is the foundation everything else builds on.